

Design Specification

Memory Management Tools

Anvil Memory

Author: UniquesKernel

Date: August 10, 2025

Contents

1	Introduction to the Problem	2
2	Memory Management	3
2.1	Scope-Centralized Memory Management	3
2.2	Initialization-Centralized Memory Management	4
3	Allocation Strategy	6

1 Introduction to the Problem

Consider a program $P \equiv (I, M)$ where I denotes an ordered sequence of instructions and M denotes an ordered sequence of available memory addresses. Let $A \subset I$ denote the set of allocation instructions and let $B \subset I$ denote the set of deallocation instructions. For P to be temporally memory safe, it must satisfy that for all execution paths $\pi \in \Pi(P)$ there exists a bijective mapping $T_\pi : A_\pi \rightarrow B_\pi$. Moreover P exhibits exponential complexity, with a lower bound of $2^{|C|}$ where $C \subset I$ denotes the set of conditional instructions.

Individual tracking of allocation-deallocation pairs admits no systematic method for managing this complexity, as programs with merely $|C| > 10$ conditional instructions yield $2^{|C|} > 1024$ execution paths. This is because such methods operate at the incorrect granularity. Furthermore, traditional approaches such as RAI and reference counting, while offering systematic tracking mechanisms, operate at the same granularity. These methods merely elevate individual allocation management to a higher level of abstraction without reducing the inherent complexity.

We reduce this complexity by observing that instructions in I exhibit temporal dependencies that admit natural groupings, termed **scopes**. Scopes may be nested, creating hierarchical structure between allocation-deallocation pairs. These hierarchical structures cause spatial fragmentation, separating allocation and deallocation pairs in I making verification of the bijective mapping error prone. By managing all allocations within a scope as a unified set, verification complexity reduces from exponential in C to linear verification in the root scope. Since nested scope allocations belong to the root set, destroying the set eliminates all nested allocations, regardless of execution path.

2 Memory Management

Let a program $P \equiv (i_n)_{n \in \mathbb{N}}$ be an ordered sequence of instructions. Let $A \subset P$ denote the set of allocation instructions and $B \subset P$ the deallocation instructions.

Definition 2.1 Memory-safe Program.

A program P is *memory-safe* if for all execution paths $\pi \in \Pi(P)$:

1. Exists bijection $T_\pi : A_\pi \rightarrow B_\pi$ where:
 - $\forall a_i \in A_\pi, T_\pi(a_i) = b_i$ is unique.
 - a_i lexically precedes b_i on π .
2. The deallocation graph $G_\pi = (V_\pi, E_\pi)$ is acyclic.
3. $\text{mem}(a_i) \cap \text{mem}(a_j) = \emptyset$ for $i \neq j$.

2.1 Scope-Centralized Memory Management

Definition 2.2 Memory Arena.

A *memory arena* $\mathcal{A} \equiv (\mathcal{M}, F, \Lambda)$ where:

- $\mathcal{M} = \{\mathcal{M}_1, \dots, \mathcal{M}_k\}$ - is a finite set of disjoint memory blocks..
- $F \subseteq \bigcup_{i=1}^k \mathcal{M}_i$ - is the set of free addresses.
- Λ - is an allocator with allocation strategy σ .

The *allocation function* $H : \mathcal{A} \rightarrow \mathcal{A} \times \text{addr}(\mathcal{M})$ is:

$$H((\mathcal{M}, F, \Lambda)) = \begin{cases} ((\mathcal{M}, F \setminus \{a\}, \Lambda), a) & \text{if } a \in F \text{ via } \sigma \\ \text{undefined} & \text{otherwise} \end{cases}$$

Subject to invariants:

1. Disjointness: $\mathcal{M}_i \cap \mathcal{M}_j = \emptyset \forall i \neq j$.
2. No individual deallocations: $\nexists b_i \in B$ for $a_i \in \mathcal{A}$.
3. Mass reclamation: $\bigcup \mathcal{M}_\pi \subseteq F_\pi^{\text{final}}$ at scope exit.

For control-flow graph G with scope entry/exit nodes G_a, G_b :

- Arena instantiation at G_a creates $\mathcal{A} = (\mathcal{M}, F, \Lambda)$.
- All paths $S_{a \rightarrow b}$ allocate via $H((\mathcal{M}, F, \Lambda))$.
- Destruction at G_b enforces $F^{\text{final}} = \bigcup \mathcal{M}$, satisfying Definition 2.1.

This formalism reduces verification complexity from $O(2^{|C|})$ paths to $O(1)$ scope-level invariants. The allocator Λ implements σ through H while maintaining:

- $$\forall \pi \in \Pi(P), 1. T_\pi \text{ bijection} \Rightarrow F^{\text{final}} = \bigcup \mathcal{M}$$
2. Lexical precedence \Rightarrow Scope nesting
 3. Disjointness $\Rightarrow \mathcal{M}_i \cap \mathcal{M}_j = \emptyset$

Memory Arena




Figure 1: Memory arena structure showing sequential allocations within contiguous blocks.

2.2 Initialization-Centralized Memory Management

We now present a further centralization of memory management through what we shall term the "Defer Pattern," which couples allocation and deallocation at the point of initialization.

Utilizing compiler facilities such as the cleanup attribute in GCC, one may specify:

```
MemoryArena *arena DEFER(memory_arena_destroy) = memory_arena_create(...);
```

This construction guarantees deallocation at scope exit, irrespective of the control flow path that leads to said exit.

The Defer Pattern exhibits the following properties:

1. **Spatial Coupling:** The allocation and its corresponding deallocation are textually adjacent in the source code, providing immediate verification of the appropriate $T_\pi(a_i) = b_i$ relationship.
2. **Compiler-Enforced Temporal Sequencing:** The compiler ensures that deallocation occurs at scope exit, enforcing the lexical ordering requirement that a_i precedes b_i .
3. **Single-Point Specification:** Memory management decisions are made once, at the point of initialization, rather than distributed throughout the control flow graph.

This approach effectively centralizes memory management responsibilities at initialization points, delegating enforcement to the compiler. The programmer need only ensure proper declaration of cleanup behaviors at allocation sites, substantially reducing the cognitive burden of tracking allocation-deallocation pairs across the program's control flow structure.

We thus observe a progression from decentralized path-based verification, to scope-centralized arena management, to initialization-centralized specification—each stage reducing the cognitive distance between allocation and deallocation logic, and consequently reducing the opportunity for memory management errors.

For compiler-enforced cleanup:

$$\underbrace{\text{MemoryArena}^* \mathcal{A}}_{\text{Allocation}} \xrightarrow[\text{Deallocation}]{\text{DEFER}} \mathcal{A} \mapsto (\mathcal{M}, F^{\text{final}}, \Lambda)_{\text{Scope exit}}$$

where the DEFER macro enforces:

$$\mathcal{A} \mapsto (\emptyset, \bigcup \mathcal{M}, \Lambda) \forall \pi \in \Pi(P).$$

3 Allocation Strategy

Definition 3.1 contiguous Memory.

Let \mathcal{M} be the set of memory addresses. Let F be the set of free addresses. A subset C of \mathcal{M} is said to be contiguous if $\exists x, y \in \mathcal{M}, C = \{a | x \leq a \leq y, x \leq y\} \wedge C \subseteq F$.

Definition 3.2 Allocation Strategy.

Let \mathcal{M} be the set of memory addresses. Let $F \subseteq \mathcal{M}$ be the set of free memory addresses. Let $A \subseteq \mathcal{M}$ be the set of allocated Memory addresses. Let $n \in \mathbb{N}$ be some natural number greater than zero. Then an allocation strategy $\sigma(F, n)$ is defined as:

$$\sigma(F, n) \rightarrow \phi((F \setminus C, A \cup C, C)),$$

Where $C \subseteq F$ is some contiguous piece of memory (see 3.1) such that $|C| = n$ and that satisfy some selection policy ϕ^a (e.g., first-fit, best-fit, worst-fit) .

^aWe shall intentionally leave ϕ as abstractly defined as it is better discussed as an implementation detail.

References